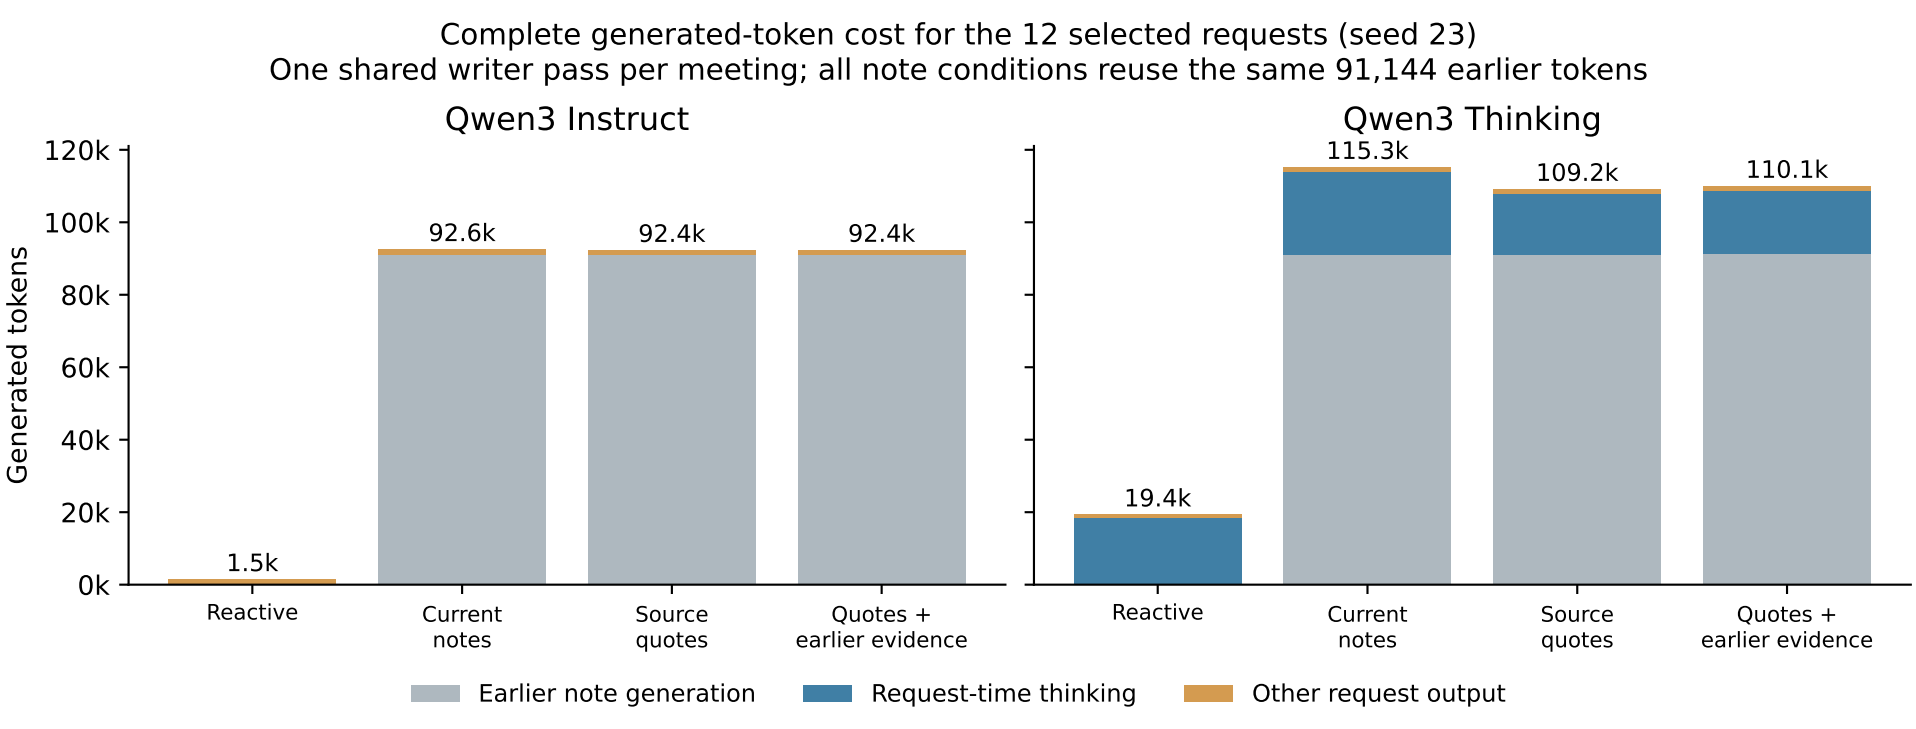
Complete generated-token cost for the 12 selected requests (seed 23)
One shared writer pass per meeting; all note conditions reuse the same 91,144 earlier tokens
Qwen3 Instruct
Qwen3 Thinking
Generated tokens
120k
100k
80k
60k
40k
20k
0k
1.5k
92.6k
92.4k
92.4k
19.4k
115.3k
109.2k
110.1k
Reactive
Current notes
Source quotes
Quotes + earlier evidence
Earlier note generation
Request-time thinking
Other request output